

# Hadoop and MapReduce


Dominic Duggan  
Stevens Institute of Technology  
Based on materials by  
A. Haeberlen, Z. Ives, Jimmy Lin

1

## **HADOOP**

2

## Recap: MapReduce dataflow



© 2013 A. Haeberlen, Z. Ives

3

## What do we need to write?

- A mapper
  - Accepts (key,value) pairs from the input
  - Produces intermediate (key,value) pairs
  - Intermediate pairs shuffled
- A reducer
  - Accepts intermediate (key,value) pairs
  - Produces final (key,value) pairs for the output
- A driver
  - Specifies which inputs to use, where to put the outputs
  - Chooses the mapper and the reducer to use

4

## Input Formats

Format	Key Type	Value Type
TextInputFormat (Default)	File Offset	Text Line
KeyValue InputFormat	Text (up to \t)	RemainingText
SequenceFile InputFormat	User-Defined	User-Defined

5

## Output Formats

Format	Description
TextOutputFormat (default)	Key \t Value \n
SequenceFileOutputFormat	Binary Serialized keys and values
NullOutputFormat	Discards Output

6

# The Mapper

Input format  
(file offset, line)

Intermediate format  
can be freely chosen

```
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.io.*;

public class FooMapper extends Mapper<LongWritable, Text, Text, Text> {
    public void map(LongWritable key, Text value, Context context) {
        context.write(new Text("foo"), value);
    }
}
```

7

# The Reducer

Intermediate format  
(same as mapper output)

Output format

```
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.io.*;

public class FooReducer extends Reducer<Text, Text, IntWritable, Text> {
    public void reduce(Text key, Iterable<Text> values, Context context)
        throws java.io.IOException, InterruptedException {
        for (Text value: values)
            context.write(new IntWritable(4711), value);
    }
}
```

Note: We may get  
multiple values for  
the same key!

8

# The Driver

```
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class FooDriver {
    public static void main(String[] args) throws Exception {
        Job job = new Job();
        job.setJarByClass(FooDriver.class);

        FileInputFormat.addInputPath(job, new Path("in"));
        FileOutputFormat.setOutputPath(job, new Path("out"));

        job.setMapperClass(FooMapper.class);
        job.setReducerClass(FooReducer.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(Text.class);

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

Mapper&Reducer are  
in the same Jar as  
FooDriver

Input and Output  
paths

Format of the (key,value)  
pairs output by the  
reducer

9

## HDFS

10

# HDFS

- Master-Worker Architecture
- Single NameNode
- Many (Thousands) DataNodes


11

## HDFS Master (NameNode)

- Manages filesystem namespace
- File metadata (“inodes”)
- Mapping “inode” to list of blocks and locations
- Authorization and Authentication
- Checkpoint & journal namespace changes

12

# Namenode



- State stored in two files: fsimage and edits
  - fsimage: Snapshot of file system metadata
  - edits: Changes since last snapshot

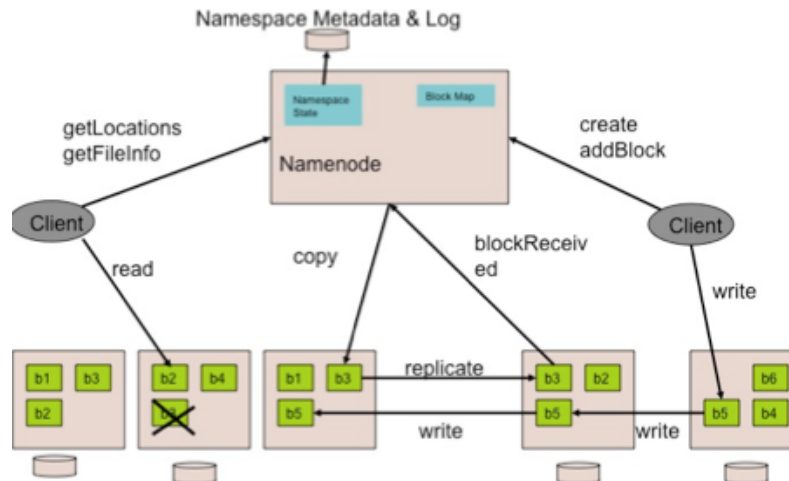
13

## Secondary Namenode

- What if the state of the namenode is lost?
- Solution #1: Metadata backups
- Solution #2: Secondary Namenode
  - Has a copy of the metadata
  - Periodically merge edit log with fsimage
  - Lags lead to data loss

14

# HDFS Architecture



15

## Accessing data in HDFS

```
$ ls -la /tmp/hadoop-ahae/dfs/data/current/
total 209588
drwxrwxr-x 2 ahae ahae 4096 2013-10-08 15:46 .
drwxrwxr-x 5 ahae ahae 4096 2013-10-08 15:39 ..
-rw-rw-r-- 1 ahae ahae 11568995 2013-10-08 15:44 blk_-3562426239750716067
-rw-rw-r-- 1 ahae ahae 90391 2013-10-08 15:44 blk_-3562426239750716067_1020.meta
-rw-rw-r-- 1 ahae ahae 4 2013-10-08 15:40 blk_5467088600876920840
-rw-rw-r-- 1 ahae ahae 11 2013-10-08 15:40 blk_5467088600876920840_1019.meta
-rw-rw-r-- 1 ahae ahae 67108864 2013-10-08 15:44 blk_7080460240917416109
-rw-rw-r-- 1 ahae ahae 524295 2013-10-08 15:44 blk_7080460240917416109_1020.meta
-rw-rw-r-- 1 ahae ahae 67108864 2013-10-08 15:44 blk_-8388309644856805769
-rw-rw-r-- 1 ahae ahae 524295 2013-10-08 15:44 blk_-8388309644856805769_1020.meta
-rw-rw-r-- 1 ahae ahae 67108864 2013-10-08 15:44 blk_-9220415087134372383
-rw-rw-r-- 1 ahae ahae 524295 2013-10-08 15:44 blk_-9220415087134372383_1020.meta
-rw-rw-r-- 1 ahae ahae 158 2013-10-08 15:40 VERSION
$
```

- HDFS implements a separate namespace
  - Only blocks and block metadata are visible

16



## Accessing data in HDFS

```
$ /usr/local/hadoop/bin/hadoop fs -ls /user/ahae
Found 4 items
-rw-r--r-- 1 ahae supergroup      1366 2013-10-08 15:46 /user/ahae/README.txt
-rw-r--r-- 1 ahae supergroup        0 2013-10-08 15:35 /user/ahae/input
-rw-r--r-- 1 ahae supergroup        0 2013-10-08 15:39 /user/ahae/input2
-rw-r--r-- 1 ahae supergroup 212895587 2013-10-08 15:44 /user/ahae/input3
$
```

- Examples:
  - `hadoop fs -put [file] [hdfsPath]` Stores a file in HDFS
  - `hadoop fs -ls [hdfsPath]` List a directory
  - `hadoop fs -get [hdfsPath] [file]` Retrieves a file from HDFS
  - `hadoop fs -rm [hdfsPath]` Deletes a file in HDFS
  - `hadoop fs -mkdir [hdfsPath]` Makes a directory in HDFS

17

## HDFS Java API

```
// Get default file system instance
fs = Filesystem.get(new Configuration());

// Or Get file system instance from URI
fs = Filesystem.get(URI.create(uri),
    new Configuration());

// Create, open, list, ...
OutputStream out = fs.create(path, ...);
InputStream in = fs.open(path, ...);
boolean isDone = fs.delete(path, recursive);
FileStatus[] fstat = fs.listStatus(path);
```

18

## INSTALLING AND RUNNING

19

### Prerequisites for Hadoop

- Java 1.6+
  - Recommended: create a [hadoop](#) user
    - Standard privileges only
  - SSH
    - Need ssh key for access to hosts in cluster
- ```
$ ssh-keygen -t rsa -P ""  
$ cat $HOME/.ssh/id_rsa.pub >> \  
$HOME/.ssh/authorized_keys
```

20

## Download Hadoop

- Download from Apache:
  - <http://hadoop.apache.org/releases.html#Download>
- Install (e.g. in /usr/local)

```
$ cd /usr/local
$ gunzip $HOME/Downloads/hadoop-2.X.X.tar.gz
$ sudo tar xvf $HOME/Downloads/hadoop-2.X.X.tar
$ sudo ln -s hadoop-2.X.X hadoop
$ sudo chown -R hadoop:staff hadoop-2.X.X hadoop
```

21

## Setting Your User Profile

- Edit `$HOME/.bash_profile`
- Add these lines:

```
export HADOOP_PREFIX="/usr/local/hadoop"
export HADOOP_HOME="${HADOOP_PREFIX}"
export HADOOP_COMMON_HOME="${HADOOP_PREFIX}"
export HADOOP_CONF_DIR="${HADOOP_PREFIX}/etc/hadoop"
export HADOOP_HDFS_HOME="${HADOOP_PREFIX}"
export HADOOP_MAPRED_HOME="${HADOOP_PREFIX}"
export HADOOP_YARN_HOME="${HADOOP_PREFIX}"
export "PATH=${PATH}:${HADOOP_PREFIX}/bin:
        ${HADOOP_PREFIX}/sbin"
```
- Load into current shell

```
. $HOME/.bash_profile
```

22

## Configuring Hadoop: `hadoop-env.sh`

- Located in `$HADOOP/etc/hadoop/hadoop-env.sh`
- Edit to set `JAVA_HOME`
  - E.g. on MacOS:  
`export JAVA_HOME=`/usr/libexec/java_home -v 1.7+``

23

## Manual compilation

- Step #1: Put `hadoop-core-...jar` into classpath:  
`export CLASSPATH=$CLASSPATH:/path/to/hadoop/hadoop-core-...jar`
- Step #2: Compile mapper, reducer, driver:  
`javac FooMapper.java FooReducer.java FooDriver.java`
- Step #3: Package into a JAR file:  
`jar cvf Foo.jar *.class`
- Alternative: "Export..." / "Java JAR file" in Eclipse

24

## Standalone Operation

- One Java process

```
$ mkdir input
```

```
$ cp ... input
```

```
$ hadoop jar jarfile-name class-name \  
    input output program-arguments
```

```
$ cat output/*
```

25

## Pseudo-Distributed Operation

- One site, one Java process per node
- `$HADOOP_HOME/etc/hadoop/core-site.xml`

```
<configuration>  
  <property>  
    <name>fs.defaultFS</name>  
    <value>hdfs://localhost:9000</value>  
  </property>  
</configuration>
```

26

## Pseudo-Distributed Operation

- One site, one Java process per node
- `$HADOOP_HOME/etc/hadoop/hdfs-site.xml`

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
</configuration>
```

27

## Pseudo-Distributed Operation

- Format the file system

```
$ hdfs namenode -format
```
- Start daemons (NameNode and DataNode)

```
$ start-dfs.sh
```

  - Log written to `$HADOOP_HOME/logs`
  - REST interface for NameNode: <http://localhost:50070>
- Make HDFS directories

```
$ hdfs dfs -mkdir /user
$ hdfs dfs -mkdir /user/joe
```

28

## Pseudo-Distributed Operation

- Copy input files into the Hadoop file system  
`$ hdfs dfs -put ... input`
- Run examples  
`$ hadoop jar jarfile class-name \  
input output arguments`
- Examine output  
`$ hdfs dfs -get output output  
$ cat output/*`
- Stop the daemons  
`$ stop-dfs.sh`

29

## Pseudo-Distributed with Yarn

- etc/hadoop/mapred-site.xml:  

```
<configuration>  
  <property>  
    <name>mapreduce.framework.name</name>  
    <value>yarn</value>  
  </property>  
</configuration>
```
- etc/hadoop/yarn-site.xml:  

```
<configuration>  
  <property>  
    <name>yarn.nodemanager.aux-services</name>  
    <value>mapreduce_shuffle</value>  
  </property>  
</configuration>
```

30

## Pseudo-Distributed with Yarn

- Start ResourceManager and NodeManager daemons:

```
$ start-yarn.sh
```

- REST interface for ResourceManager:

<http://localhost:8088>

- Run a MR job
- Stop the daemons:

```
$ stop-yarn.sh
```

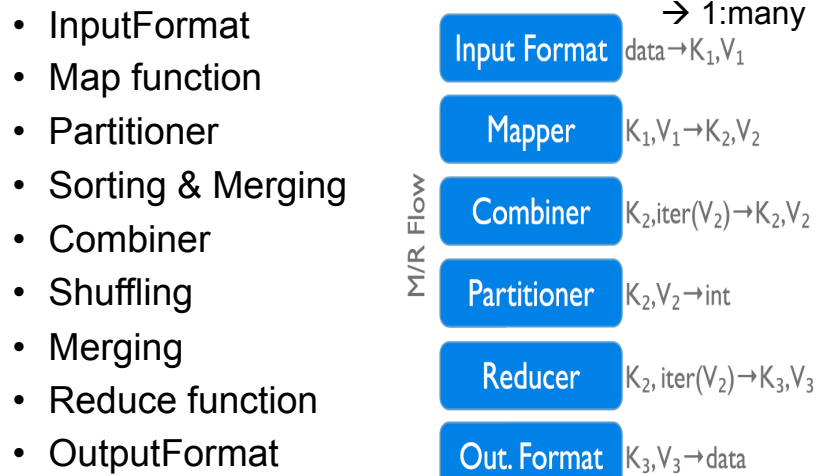
31

## HADOOP COMPONENTS

32

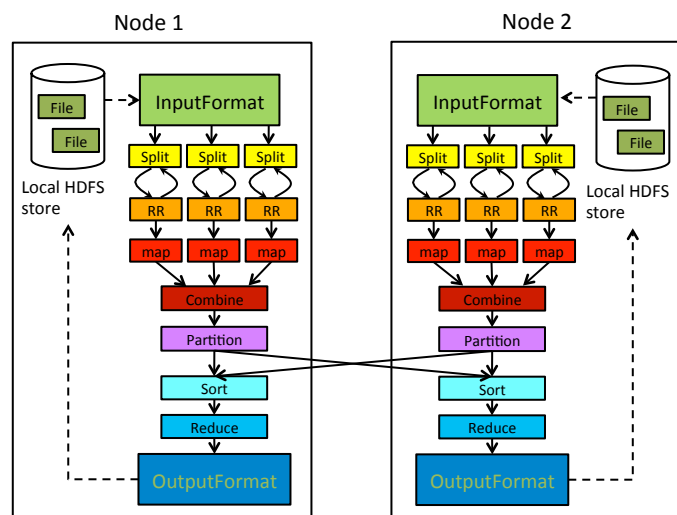


# Data Flow in Hadoop



33

## Detailed dataflow in Hadoop




© 2013 A. Haeberlen, Z. Ives

34

# Input Format


- Input files and format
  - Defaults provided, e.g., TextInputFormat, DBInputFormat, KeyValueTextInputFormat...
- Defines InputSplits
  - Break file into separate tasks
  - Example: one task for each 64MB block
- Factory for RecordReaders
  - RecordReaders read the file into (key,value) pairs
  - TextInputFormat: byte offset in file as key
  - KeyValueInputFormat: key is everything up to first tab



35

# Combiners


- Optional component after mappers
  - Input: All data emitted by mappers on a given node
  - Output passed to the partitioner
- Why is this useful?
  - Word count emits (xyz, 1) pairs for each word xyz
  - Pass (xyz, k) to the reducer



36

## Partitioner


- Which intermediate key-value pairs should go to which reducer
- Defines a partition on the set of KV pairs
  - Number of partitions = number of reducers
- Default partitioner (HashPartitioner): partition based on a hash of the key



37

## Output Format

- Counterpart to InputFormat
- Where output is stored
  - Factory for RecordWriter
- Several implementations provided
  - TextOutputFormat (default)
  - DBOutputFormat
  - MultipleTextOutputFormat
  - ...



38

# Hadoop daemons

- TaskTracker
  - Runs maps and reduces. One per node.
- JobTracker
  - Accepts jobs; assigns tasks to TaskTrackers
- DataNode
  - Stores HDFS blocks
- NameNode
  - Stores HDFS metadata
- SecondaryNameNode
  - Merges edits file with snapshot; "backup" for NameNode


A single node can run  
more than one of these!

39


## An example configuration



Small cluster



Medium cluster



© 2013 A. Haeberlen, Z. Ives


40

## Fault tolerance

- What if a node fails during a job?
  - JobTracker re-executes the failed node's tasks
- What specifically should be re-executed?
  - Depends on the phase the job was in
  - Mapping phase: Re-execute all maps assigned to failed node
  - Reduce phase: Re-execute all reduces assigned to node
    - Need to re-execute map tasks on the failed node as well!

41

## Placement and locality



- Which of the replicated blocks should be read?
  - If possible, pick the closest one (reduces network load)
  - Distance metric takes into account: Nodes, racks, datacenters
- Where should the replicas be put?
  - Tradeoff between fault tolerance and locality/performance

42

## GRAPH ALGORITHMS


43

### Beyond average/sum/count

- Networks of relationships and shared features
  - Members of a social network
  - Customers
  - The Web (documents with links)
  - Documents: topics, words, authors, etc.

44


## Thinking about related objects



- Represent related objects as labeled, directed graph
- Entities == nodes
  - Nodes: IDs
- Relationships == edges
  - Edges: values

45

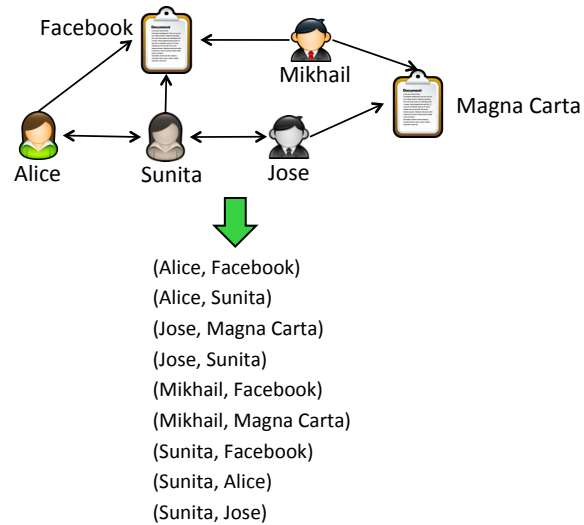
## Encoding the data in a graph



- $G = (V, E)$  where  $V$  is vertices,  $E$  is edges of the form  $(v_1, v_2)$  where  $v_1, v_2 \in V$
- Assume we only care about connected vertices
  - Then we can capture a graph simply as the **edges**
  - ... or as an **adjacency list**:  $v_i$  goes to  $[v_j, v_{j+1}, \dots]$

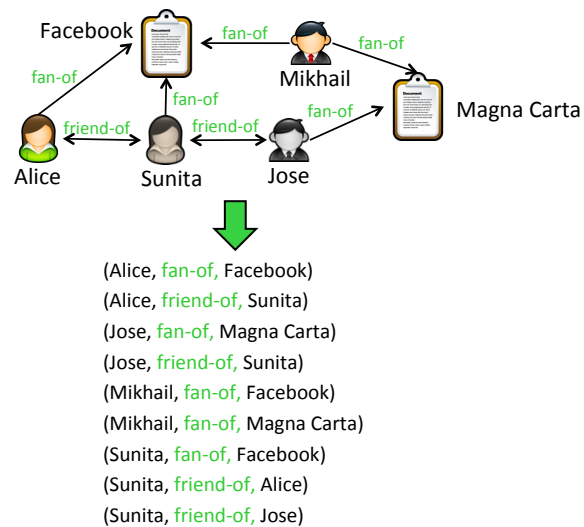
46

## Graph encodings: Set of edges



47

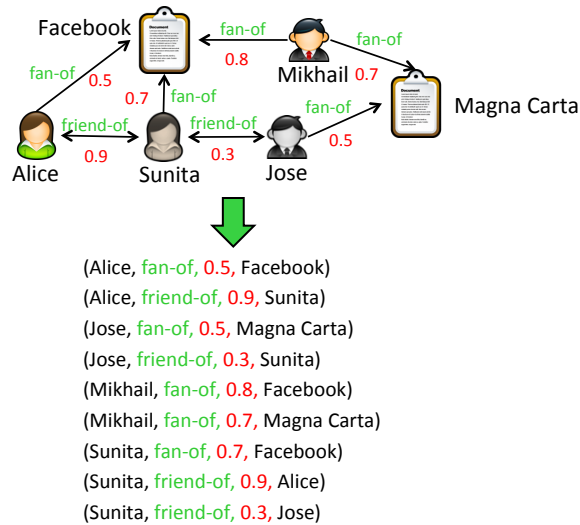
## Graph encodings: Adding edge types



48




## Graph encodings: Adding weights



49

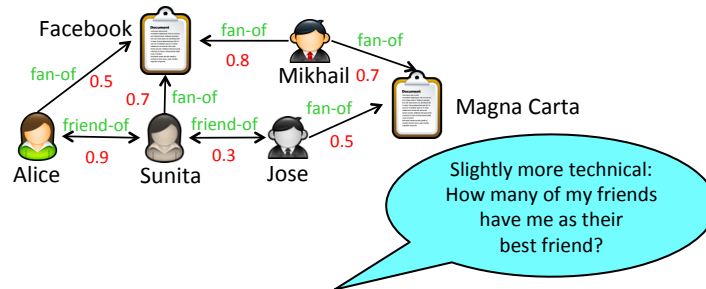
## A computation model for graphs



- Perform computations
  - Simple example: Which users are their friends' best friend?
- Method
  - annotating the vertices with additional information
  - propagating the information along the edges

50


## A computation model for graphs



- Example: Am I my friends' best friend?
  - Step #1: Discard irrelevant vertices and edges

51

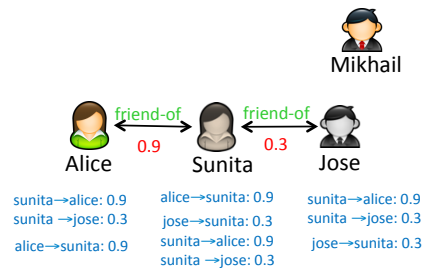
## A computation model for graphs



- Example: Am I my friends' best friend?
  - Step #1: Discard irrelevant vertices and edges
  - Step #2: Annotate each vertex with list of friends

52

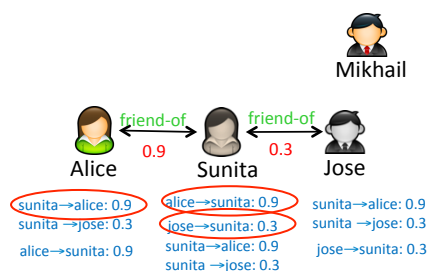
## A computation model for graphs



- Example: Am I my friends' best friend?
  - Step #1: Discard irrelevant vertices and edges
  - Step #2: Annotate each vertex with list of friends
  - Step #3: Push annotations along each edge

53

## A computation model for graphs



- Example: Am I my friends' best friend?
  - Step #1: Discard irrelevant vertices and edges
  - Step #2: Annotate each vertex with list of friends
  - Step #3: Push annotations along each edge
  - Step #4: Determine result at each vertex

54

## Can we do this in MapReduce?

```
map(key: node, value: list of <otherNode, relType, strength>)
{
}
reduce(key: _____, values: list of _____)
{
}
```

- Using adjacency list representation?

55

## Can we do this in MapReduce?

```
map(key: node, value: <otherNode, relType, strength>)
{
}
reduce(key: _____, values: list of _____)
{
}
```

- Using single-edge data representation?

56

## Generalizing...

- Beyond direct friend relationships
  - Example: How many of my friends' friends (distance-2 neighbors) have me as their best friend's best friend?
- How about distance  $k > 2$ ?
- Requires multiple iterations of MapReduce!

57

## Iterative MapReduce

```
copy files from input dir → staging dir 1
(optional: do some preprocessing)

while (!terminating condition) {
  map from staging dir 1
  reduce into staging dir 2
  move files from staging dir 2 → staging dir 1
}

(optional: postprocessing)
move files from staging dir 2 → output dir
```

- Reduce output must be compatible with the map input
  - What can happen if we filter out some information in the mapper or in the reducer?

58

## Graph algorithms and MapReduce

- Multiple map/reduce stages processing one “wave” at a time
  - Iterative MapReduce
  - Chains of map/reduce

59

## PATH-BASED ALGORITHMS

60


## Path-based algorithms

- Compute information about the paths (sets of paths) between nodes
  - Edges may be annotated with cost, distance, or similarity
- Examples:
  - Shortest path from one node to another
  - Minimum spanning tree (minimal-cost tree connecting all vertices in a graph)
  - Steiner tree (minimal-cost tree connecting certain nodes)
  - Topological sort (node in a DAG comes before all nodes it points to)

61

## Single-Source Shortest Path (SSSP)

- Given directed graph where each edge  $(u,v)$  has cost  $\text{dist}(u,v)$
- Given a start node  $s$
- Compute min cost path from  $s$  to each other node



62

## SSSP: Intuition

- The shortest path follows the **principle of optimality**: the last step  $(u,v)$  makes use of the shortest path to  $u$

```
bestDistance(v) {  
  if (v == source) {  
    return distance 0  
  } else {  
    du = min{ bestDistance(u) + dist(u,v) | u adjacent to v}  
    return du  
  }  
}
```

63

## SSSP: Intuition

- The shortest path follows the **principle of optimality**: the last step  $(u,v)$  makes use of the shortest path to  $u$

```
bestDistanceAndPath(v) {  
  if (v == source) {  
    return (distance 0, list(v))  
  } else {  
    du = min{ d + dist(u,v) | u adjacent to v and  
              bestDistanceAndPath(u) == (d, L) }  
    u = node with bestDistanceAndPath(u) == (d, _)  
        and d + dist(u,v) == du  
    (d,L) = bestDistanceAndPath(u)  
    return (distance du, list L with v added)  
  }  
}
```

64



## SSSP: Solution

- Traditional approach: **Dijkstra's algorithm**

```

V: vertices, E: edges, S: start node


foreach v in V
    dist_S_to[v] = infinity
    predecessor[v] = nil
spSet = {}
Q := V
while (Q not empty) do
    u := Q.removeNodeClosestTo(S)
    spSet := spSet + {u}
    foreach v in V where (u,v) in E
        if (dist_S_to[v] > dist_S_to[u] + dist(u,v)) {
            dist_S_to[v] = dist_S_to[u] + dist(u,v)
            predecessor[v] = u
        }
    
```

Initialize length and  
last step of path  
to default values

Update length and  
path based on edges  
radiating from u

65

## SSSP: Dijkstra in Action




```

Q = {s,a,b,c,d}      spSet = {}
dist_S_to: {(a,infinity), (b,infinity), (c,infinity), (d,infinity)}
predecessor: {(a,nil), (b,nil), (c,nil), (d,nil)}
    
```

66

Example from CLRS 2nd ed. p. 528

## SSSP: Dijkstra in Action




$Q = \{a, b, c, d\}$        $spSet = \{s\}$   
 $dist\_S\_To: \{(a, 10), (b, \infty), (c, 5), (d, \infty)\}$   
 $predecessor: \{(a, s), (b, nil), (c, s), (d, nil)\}$

67

Example from CLRS 2nd ed. p. 528

## SSSP: Dijkstra in Action




$Q = \{a, b, d\}$        $spSet = \{c, s\}$   
 $dist\_S\_To: \{(a, 8), (b, 14), (c, 5), (d, 7)\}$   
 $predecessor: \{(a, c), (b, c), (c, s), (d, c)\}$

68

Example from CLRS 2nd ed. p. 528

## SSSP: Dijkstra in Action




$Q = \{a, b\}$        $spSet = \{c, d, s\}$   
 $dist\_S\_To: \{(a, 8), (b, 13), (c, 5), (d, 7)\}$   
 $predecessor: \{(a, c), (b, d), (c, s), (d, c)\}$

69

Example from CLRS 2nd ed. p. 528

## SSSP: Dijkstra in Action




$Q = \{b\}$        $spSet = \{a, c, d, s\}$   
 $dist\_S\_To: \{(a, 8), (b, 9), (c, 5), (d, 7)\}$   
 $predecessor: \{(a, c), (b, a), (c, s), (d, c)\}$

70

Example from CLRS 2nd ed. p. 528

## SSSP: Dijkstra in Action




$Q = \{\}$                        $spSet = \{a, b, c, d, s\}$   
 $dist\_S\_To: \{(a, 8), (b, 9), (c, 5), (d, 7)\}$   
 $predecessor: \{(a, c), (b, a), (c, s), (d, c)\}$

71

## SSSP: How to parallelize?

- Dijkstra single route at a time
  - No real parallelism
- Alternatively:
  - “radiate” from the origin
  - one “edge hop distance” at a time



72


## SSSP: Revisiting the inductive definition

```
bestDistance(v) {  
  if (v == source) {  
    return distance 0  
  } else {  
    du = min{ bestDistance(u) + dist(u,v) | u adjacent to v}  
    return du  
  }  
}
```

- Dijkstra's algorithm
  - Select min u (**prunes** certain points)
- Instead look at all potential **u**'s
  - Compute iteratively
  - “frontier set” of **u** nodes

73

## SSSP: MapReduce formulation



Map Input:  
 $x \rightarrow \dots\{<z, d1>\}$   
 $y \rightarrow \dots\{<z, d2>\}$

Map Output:  
 $z \rightarrow \{<x, dx+d1>\}$   
 $z \rightarrow \{<y, dy+d2>\}$

Reduce: pick min of  
 $\{dx+d1, dy+d2\}$

Reduce Output:  
 $z \rightarrow dz, x \text{ or } y, \dots$

74

# SSSP: MapReduce formulation

- init:
  - For each node,  $\text{node ID} \rightarrow \langle \infty, -, \{\langle \text{succ-node-ID}, \text{edge-cost} \rangle\} \rangle$ 

The shortest path we have found so far from the source to nodeID has length  $\infty$ ...
- map:
  - take  $\text{node ID} \rightarrow \langle \text{distance}, \text{next}, \{\langle \text{succ-node-ID}, \text{edge-cost} \rangle\} \rangle$
  - For each succ-node-ID:
    - emit  $\text{succ-node ID} \rightarrow \langle \text{node ID}, \text{distance} + \text{edge-cost} \rangle$ 


This is a new path from the source to succ-node-ID that we just discovered (not necessarily shortest)
    - emit  $\text{node ID} \rightarrow \langle \text{distance}, \text{next}, \{\langle \text{succ-node-ID}, \text{edge-cost} \rangle\} \rangle$ 

Why is this necessary?
- reduce:
  - distance := min cost from a predecessor; next := that predec.
  - emit  $\text{node ID} \rightarrow \langle \text{distance}, \text{next}, \{\langle \text{succ-node-ID}, \text{edge-cost} \rangle\} \rangle$
- Repeat until no changes
- Postprocessing: Remove adjacency lists

75

## Iteration 0: Base case


mapper:  $(a, \langle s, 10 \rangle) (c, \langle s, 5 \rangle)$  edges  
 reducer:  $(a, \langle 10, \dots \rangle) (c, \langle 5, \dots \rangle)$



76

## Iteration 1


mapper: (a,<s,10>) (c,<s,5>) (a,<c,8>) (c,<a,9>) (b,<a,11>)  
 (b,<c,14>) (d,<c,7>) edges  
 reducer: (a,<8, ...>) (c,<5, ...>) (b,<11, ...>) (d,<7, ...>)



77

## Iteration 2

mapper: (a,<s,10>) (c,<s,5>) (a,<c,8>) (c,<a,9>) (b,<a,11>) (b,<c,14>) (d,<c,7>)  
 (b,<d,13>) (d,<b,15>) edges  
 reducer: (a,<8>) (c,<5>) (b,<11>) (d,<7>)




78

## Iteration 3

No change!  
Convergence!

mapper: (a,<s,10>) (c,<s,5>) (a,<c,8>) (c,<a,9>) (b,<a,11>)  
(b,<c,14>) (d,<c,7>) (b,<d,13>) (d,<b,15>) edges  
reducer: (a,<8>) (c,<5>) (b,<11>) (d,<7>)



Question: If a vertex's path cost is the same in two consecutive rounds, can we be sure that this vertex has converged?

79

## CLUSTERING

80




## Learning (clustering / classification)

- Group related entities
  - Clustering: based on similarity
  - Classification: based on putting them into a semantically meaningful class
- Both are instances of machine learning

81

## The k-clustering Problem



- Given: A set of items in a n-dimensional **feature space**
  - Example: data points from survey, people in a social network
- Goal: Group the items into k “clusters”

82

## Approach: k-Means

- Let  $m_1, m_2, \dots, m_k$  be representative points for each of our  $k$  clusters
  - Specifically: the **centroid** of the cluster
- Initialize  $m_1, m_2, \dots, m_k$  to random values in the data
- For  $t = 1, 2, \dots$ :
  - Map each observation to the closest mean


$$S_i^{(t)} = \{x_j : \|x_j - m_i^{(t)}\| \leq \|x_j - m_{i^*}^{(t)}\|, i^* = 1, \dots, k\}$$

- Assign the  $m_i$  to be a new centroid for each set

$$m_i^{(t+1)} = \frac{1}{|S_i^{(t)}|} \sum_{x_j \in S_i^{(t)}} x_j$$


83

## A simple example (1/4)




84

## A simple example (2/4)




85

## A simple example (3/4)



86

## A simple example (4/4)



87

## k-Means in MapReduce

- Map #1:
  - Input:  
node ID  $\rightarrow$  <position, centroid ID, {centroid IDs and positions}>
  - Compute nearest centroid; emit  
centroid ID  $\rightarrow$  <node ID, position>
- Reduce #1:
  - Recompute centroid position from positions of nodes in it
  - Emit centroidID  $\rightarrow$  <node IDs, positions> and for all other centroid IDs, emit  
otherCentroidID  $\rightarrow$  centroid(centroidID, X, Y)
    - Each centroid will need to know where all the other centroids are

88

- Map #2:
  - Pass through values to Reducer #2
- Reduce #2:
  - For each node in the current centroid, emit  
node ID → <position, centroid ID,  
{centroid IDs and positions}>
    - Input for the next map iteration
  - Also, emit <X, <centroid ID, position>>
    - This will be the 'result'
- Repeat until no change

## CLASSIFICATION

45

# Classification



- Suppose we want to learn what is spam (or interesting, or ...)
  - Predefine a set of **classes** with semantic meaning
  - **Train** an algorithm to look at data and assign a class
    - Based on giving it some **examples** of data in each class
    - ... and the sets of **features** they have

91

## A simple example

- Look at the keywords in the email's title:

Message(1, "Won contract")  
 Message(2, "Won award")  
 Message(3, "Won the lottery")  
 Message(4, "Unsubscribe")  
 Message(5, "Millions of customers")  
 Message(6, "Millions of dollars")

- What is **probability** message "Won Millions" is ?

$$\begin{aligned}
 & p(\text{spam} | \text{containsWon}, \text{containsMillions}) \\
 &= \frac{p(\text{spam}) p(\text{containsWon}, \text{containsMillions} | \text{spam})}{p(\text{containsWon}, \text{containsMillions})}
 \end{aligned}$$

Bayes' Theorem

92

## Classification using Naïve Bayes

- Basic assumption: Probabilities of events are independent
- Under this assumption,

$$\frac{p(\text{spam}) p(\text{containsWon,containsMillions} \mid \text{spam})}{p(\text{containsWon,containsMillions})}$$

$$= \frac{p(\text{spam}) p(\text{containsWon} \mid \text{spam}) p(\text{containsMillions} \mid \text{spam})}{p(\text{containsWon}) p(\text{containsMillions})}$$

$$= 0.5 * 0.67 * 0.33 / (0.5 * 0.33) = 0.67$$

- Train a learner (compute the above probabilities) using MapReduce?

93

## Train the learner

- $p(\text{spam})$ 
  - Count how many spam emails there are Easy
  - Count total number of emails Easy
- $p(\text{containsXYZ} \mid \text{spam})$ 
  - Count how many spam emails contain XYZ 1
  - Count how many spam emails there are Easy
- $p(\text{containsXYZ})$ 
  - Count how many emails contain XYZ overall 2
  - Count total number of emails Easy

94

## Training a Naïve Bayes Learner

- map 1:
    - takes `messageId` → `<class, {words}>`
    - emits `<word, class>` → 1
  - reduce 1:
    - emits `<word, class>` → count
  - map 2:
    - takes `messageId` → `<class, {words}>`
    - emits `word` → 1
  - reduce 2:
    - emits `word` → totalCount
- Count how many emails in the class contain the word (modified WordCount)
- Count how many emails contain the word overall (WordCount)

95

## PAGERANK

96

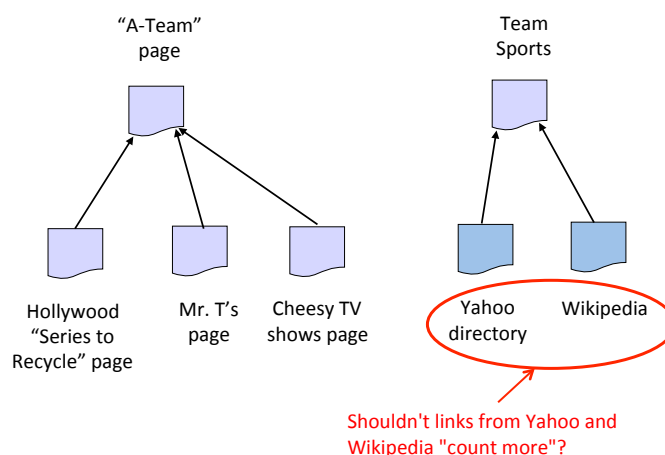


## Link analysis

- Search engine challenge
  - Problem: how to prioritize pages?
- Idea: Hyperlinks encode human judgment
  - Intra-domain links: internal navigation
  - Inter-domain links: confer authority?
- Idea: Boost the rank of pages with lots of inbound links?

97

## Problem: Popularity $\neq$ relevance!



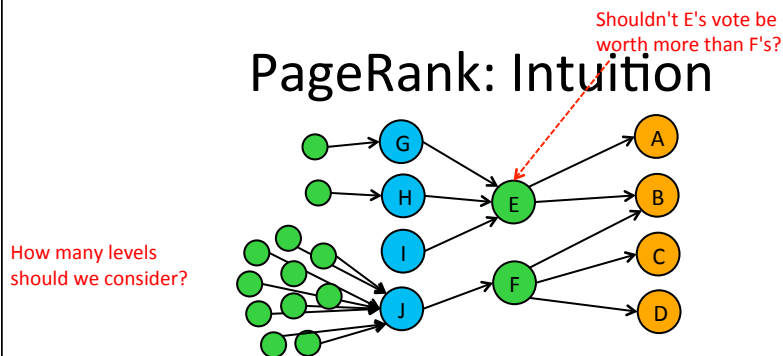
98

## Other applications

- How do we measure the "impact" of a researcher?  
(#papers? #citations?)
- Who are the most "influential" individuals in a social network?  
(#friends?)
- Which programmers are writing the "best" code?  
(#uses?)
- ...

99

## PageRank: Intuition



- Imagine a contest for The Web's Best Page
  - Initially, each page has one vote
  - A page votes for all pages it links to
  - A page's vote is split vote equally between endorsed pages
  - Voting proceeds in rounds
  - In each round, each page has the number of votes it received in the previous round

100

## PageRank

- Each page  $i$  is given a rank  $x_i$
- Goal: Assign the  $x_i$  such that the rank of each page is governed by the ranks of the pages linking to it:

$$x_i = \sum_{j \in B_i} \frac{1}{N_j} x_j$$

Rank of page  $i$       Rank of page  $j$       Number of links out from page  $j$

How do we compute the rank values?

Every page  $j$  that links to  $i$

101

## Random Surfer Model

- **Random surfer** starts on a random page and, in each step:
  - with probability  $d$ , clicks on a random link
  - with probability  $1-d$ , jumps to a random page
    - Ignoring links on page
- PageRank of a page: the fraction of steps the surfer spends on that page
  - Transition matrix can be interpreted as a Markov Chain

102

## Iterative PageRank (simplified)

Initialize all ranks to be equal, e.g.:

$$x_i^{(0)} = \frac{1}{n}$$

Iterate until convergence

$$x_i^{(k+1)} = \sum_{j \in B_i} \frac{1}{N_j} x_j^{(k)}$$


No need to decide  
how many levels  
to consider!

103

## Example: Step 0

Initialize all ranks to be equal


$$x_i^{(0)} = \frac{1}{n}$$



104

## Example: Step 1

Propagate weights  
across out-edges




105

## Example: Step 2

Compute weights  
based on in-edges


$$x_i^{(1)} = \sum_{j \in B_i} \frac{1}{N_j} x_j^{(0)}$$



106

## Example: Convergence

$$x_i^{(k+1)} = \sum_{j \in B_i} \frac{1}{N_j} x_j^{(k)}$$



107

## Naïve PageRank Algorithm Restated

- Let
  - $N(p)$  = number outgoing links from page  $p$
  - $B(p)$  = number of back-links to page  $p$

$$\text{PageRank}(p) = \sum_{b \in B(p)} \frac{1}{N(b)} \text{PageRank}(b)$$

- Each page  $b$  *distributes its importance* to all of the pages it points to (so we scale by  $1/N(b)$ )
- Page  $p$ 's importance is increased by the *importance of its back set*

108

# Linear Algebra formulation


- Create an  $m \times m$  matrix  $M$  to capture links:
  - $M(i, j) = 1 / n_j$  if page  $i$  is pointed to by page  $j$  and page  $j$  has  $n_j$  outgoing links
  - $= 0$  otherwise
  - Initialize all PageRanks to 1, multiply by  $M$  repeatedly until all values converge:

$$\begin{bmatrix} \text{PageRank}(p_1') \\ \text{PageRank}(p_2') \\ \dots \\ \text{PageRank}(p_m') \end{bmatrix} = M \begin{bmatrix} \text{PageRank}(p_1) \\ \text{PageRank}(p_2) \\ \dots \\ \text{PageRank}(p_m) \end{bmatrix}$$

- Computes **principal eigenvector** via **power iteration**

109

## Example



$$\begin{bmatrix} g' \\ y' \\ a' \end{bmatrix} = \begin{bmatrix} 0 & 0.5 & 0.5 \\ 0 & 0 & 0.5 \\ 1 & 0.5 & 0 \end{bmatrix} * \begin{bmatrix} g \\ y \\ a \end{bmatrix}$$


Running for multiple iterations:

$$\begin{bmatrix} g \\ y \\ a \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \begin{bmatrix} 1 \\ 0.5 \\ 1.5 \end{bmatrix}, \begin{bmatrix} 1 \\ 0.75 \\ 1.25 \end{bmatrix}, \dots, \begin{bmatrix} 1 \\ 0.67 \\ 1.33 \end{bmatrix}$$

Total rank sums to number of pages

110

## PageRank Sinks




Running for multiple iterations:

$$\begin{bmatrix} g \\ y \\ a \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \begin{bmatrix} 0.5 \\ 1 \\ 0.5 \end{bmatrix}, \begin{bmatrix} 0.25 \\ 0.5 \\ 0.25 \end{bmatrix}, \dots, \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

111

## PageRank Hogs




Running for multiple iterations:

$$\begin{bmatrix} g \\ y \\ a \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \begin{bmatrix} 0.5 \\ 2 \\ 0.5 \end{bmatrix}, \begin{bmatrix} 0.25 \\ 2.5 \\ 0.25 \end{bmatrix}, \dots, \begin{bmatrix} 0 \\ 3 \\ 0 \end{bmatrix}$$

112



## Stopping the Hog



Running for multiple iterations:

$$\begin{bmatrix} g \\ y \\ a \end{bmatrix} = \begin{bmatrix} 0.57 \\ 1.85 \\ 0.57 \end{bmatrix}, \begin{bmatrix} 0.39 \\ 2.21 \\ 0.39 \end{bmatrix}, \begin{bmatrix} 0.32 \\ 2.36 \\ 0.32 \end{bmatrix}, \dots, \begin{bmatrix} 0.26 \\ 2.48 \\ 0.26 \end{bmatrix}$$

... though does this seem right?

113

## Improved PageRank

- Remove out-degree 0 nodes
- Add **decay factor**  $d$  to deal with sinks

$$PageRank(p) = (1 - d) + d \sum_{b \in B_p} \frac{1}{N(b)} PageRank(b)$$

- Typical value:  $d=0.85$
- “Random surfer” Intuition:
  - Surfer occasionally stops following link sequence and jumps to new random page, with probability  $1 - d$

114

## PageRank on MapReduce

- Inputs
  - page  $\rightarrow$   $\langle \text{currentWeightOfPage}, \{\text{adjacency list}\} \rangle$
- Map
  - Page  $p$  “propagates”  $1/N_p$  of its  $d * \text{weight}(p)$  to the destinations of its out-edges (think like a vertex!)
- Reduce
  - $p$ -th page sums the incoming weights and adds  $(1-d)$ , to get its  $\text{weight}'(p)$
- Iterate until convergence

115

## PageRank on MapReduce

- Iterate until convergence
  - Common practice: run some fixed number of times, e.g., 25x
  - Alternatively: Test after each iteration with a second MapReduce job, to determine the maximum change between old and new weights

116

## Conclusions

- Common kinds of algorithms used on the Web
  - Path analysis
  - Clustering and classification
  - Link analysis
- Straightforward, often **iterative**, MapReduce formulation

117