

High Availability

Dominic Duggan

Stevens Institute of Technology

Based in part on material by Ken Birman

1

Types of reliability

- Recoverability
 - Restart in a sensible state
- High availability
 - Operational during failure
 - **Replicate critical data**

2

Replication for High Availability

- Active replication (State machine)
 - Peer-to-peer replicas
 - Each replica is **deterministic** state machine
 - **Operations** executed in same order on all replicas
 - All updates are totally ordered

3

Replication for High Availability

- Passive replication (Primary-backup)
 - Primary replica with pool of backups
 - **Operation** executed on the primary
 - **Updates** performed in same order on all replicas
- Hybrid
 - Ex: Primary-backup where operation executed on all replicas

4

Uses of replication

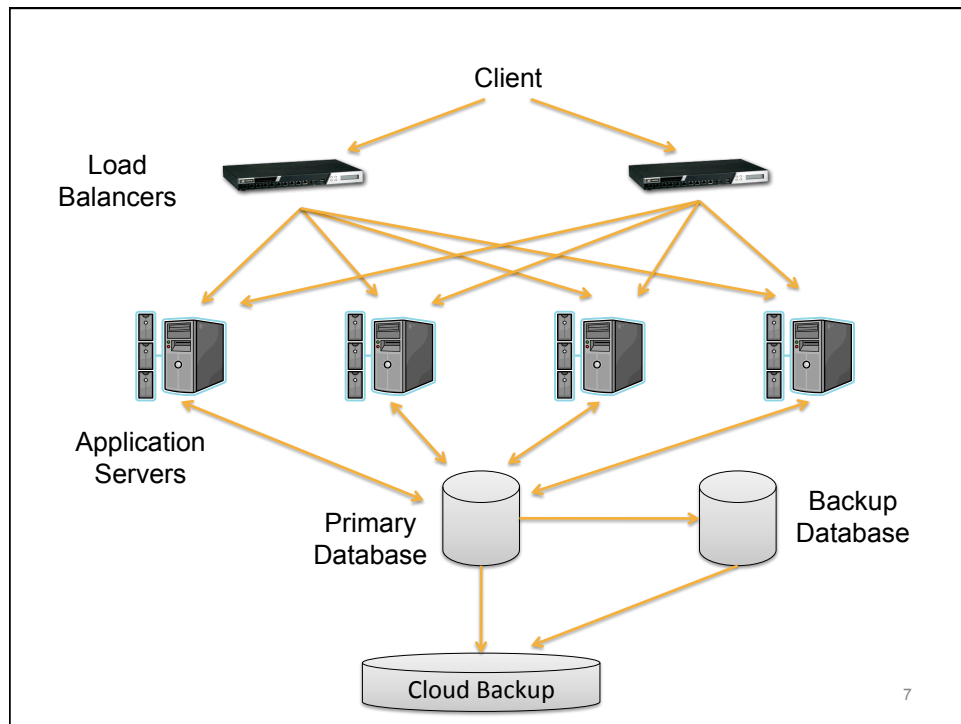
- High availability
- Share loads and improve scalability
- Replicate locking or synchronization state
- Replicate membership information in a data center (routing)
- Replicate management information to tune performance

5

Transactional Replication

- One-Copy Serializability (1SR)
 - Effect of transactions on *replicated* data items are same as if performed serially on *single* data items
 - Key: Failures and recoveries must be serialized with respect to transactions
 - Reason: Updates only performed on *available copies*


6



Server replication

- Primary sends log to backup
- Backup replays the log
 - applies committed transactions to its state
- If primary crashes, backup can take over


Primary/backup



Clients initially connected to primary, which keeps backup up to date. Backup tracks log

9


Primary/backup



Primary crashes. Backup sees the channel break.

10

Primary/backup



11

SERVER REPLICATION ISSUES


12

Issues

- Under what conditions should backup take over?
 - “Split brain” problem
- Theoretically needs 2PC to ensure that primary and backup stay in same states!


13

Split brain




14

Split brain



15

Split brain



16

Solutions

- Single server with restart
- Allow backup to “kill” the primary
 - Process groups membership service
- “Majority vote”
 - Quorum consensus




17

Issues

- Under what conditions should backup take over?
 - “Split brain” problem
- Theoretically needs 2PC to ensure that primary and backup stay in same states!


18

Primary/backup




19

Primary/backup



20

Primary/backup



21

Real systems

- Primary-backup with logging
- Omit the 2PC
 - Backup may lag state of primary
 - Hardware solutions?
 - Shared disk
 - Only one can write the disk – “token”


22

Reconciliation

- Fix transactions impacted by loss of tail of log
 - Apply the missing updates
 - Cascaded rollback
 - Worst case: human intervention
- Similar to compensations in long-lived transactions

23

Reconciliation



Reconciliation: Merge T_{k+1}, \dots, T_{k+m} into T_1', T_2', \dots while preserving data consistency

24

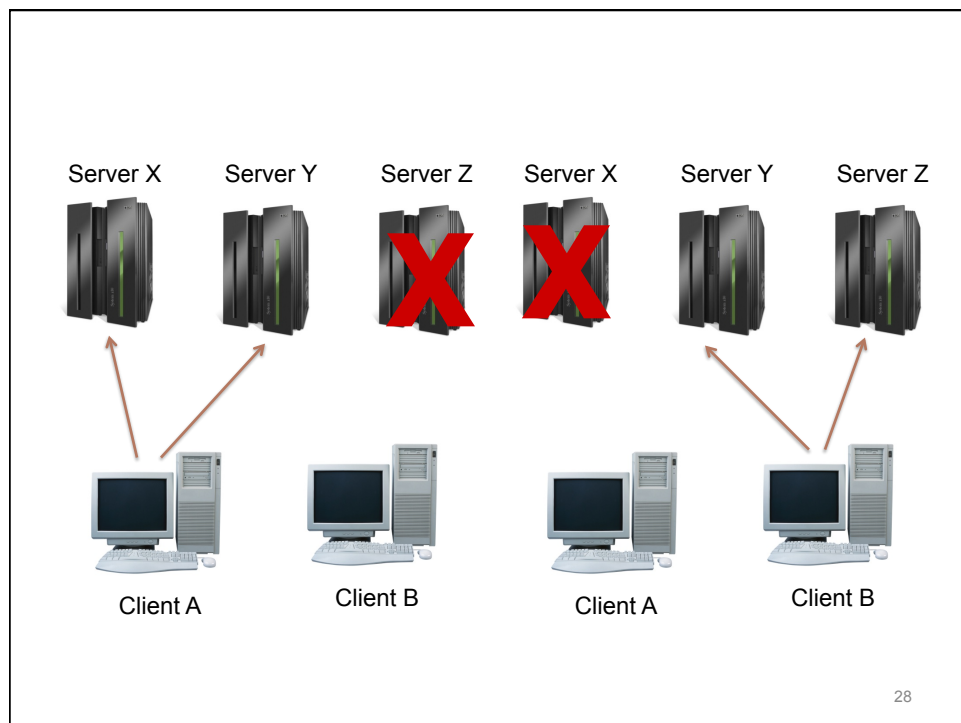
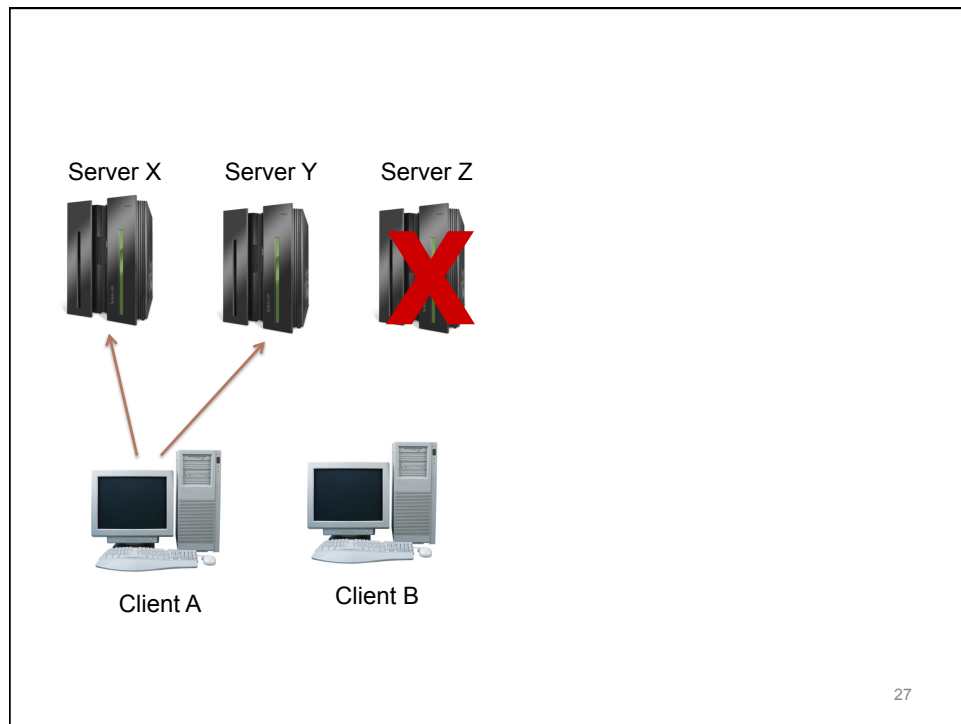
STATE MACHINE: QUORUM CONSENSUS

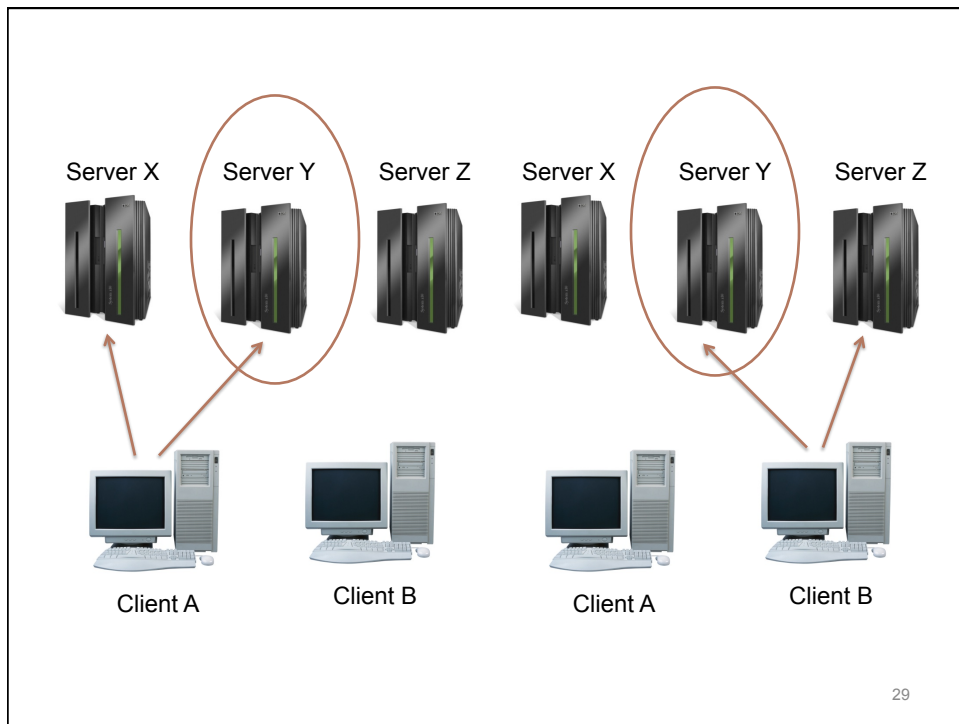
25

Issues

- How do we avoid split brain?
- How do we ensure agreement on order of updates?

26





Quorum Consensus

- Each replicated object has an update and a read quorum
- Rules
 - A quorum read should “intersect” any prior quorum write at ≥ 1 processes
 - A quorum write should also intersect any other quorum write
- So, in a group of size N :
 - $Q_r + Q_w > N$, and
 - $Q_w + Q_w > N$

30

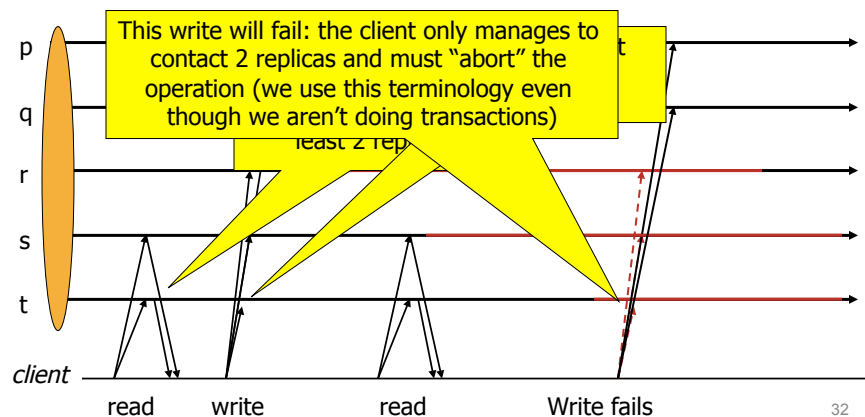
Quorum example

- X is replicated at {a,b,c,d,e}
- Possible values?
 - $Q_w = 1, Q_r = 5$ (violates $Q_w + Q_r > 5$)
 - $Q_w = 2, Q_r = 4$ (same issue)
 - $Q_w = 3, Q_r = 3$
 - $Q_w = 4, Q_r = 2$
 - $Q_w = 5, Q_r = 1$ (violates availability)
- Probably prefer $Q_w = 4, Q_r = 2$

31

Static membership example

$Q_{\text{read}} = 2, Q_{\text{write}} = 4$



32

Issues

- How do we avoid split brain?
- How do we ensure agreement on order of updates?

33

**STATE MACHINE:
ORDERING UPDATES**

34

Versions of replicated data

- Replicated data items have “versions”
 - I.e. can't just say “ $X_p=3$ ”.
 - X_p has *timestamp* $[7,q]$
 - X_p has *value* 3
 - Timestamp
 - must increase monotonically
 - includes a process id to break ties

35

Read

- Wait until Q_R processes reply
- Use value with largest timestamp
 - Break ties by looking at the pid
 - For example
 - $[6,x] < [9,a]$
 - $[7,p] < [7,q]$
 - *Even if a process owns a replica, it can't just trust its own data*

36

Write

- Can't support incremental updates
 - $x=x+1$
 - Insert into a queue
- Quorum
 - Use a commit protocol
- How to determine the version number
 - Voting protocol

37

Protocol

1. Propose the write: "I would like to set $X=3$ "
2. Members "lock" the variable against reads, *put the request into a queue of pending writes*, and send back:
"I propose time $[t, pid]$ "
Time is a logical clock.
3. Initiator collects replies, hoping to receive Q_w

$\geq Q_w$ OKs

$< Q_w$ OKs

Compute maximum of
proposed $[t, pid]$ pairs.
Commit at that time

Abort

38

Voting based on logical time


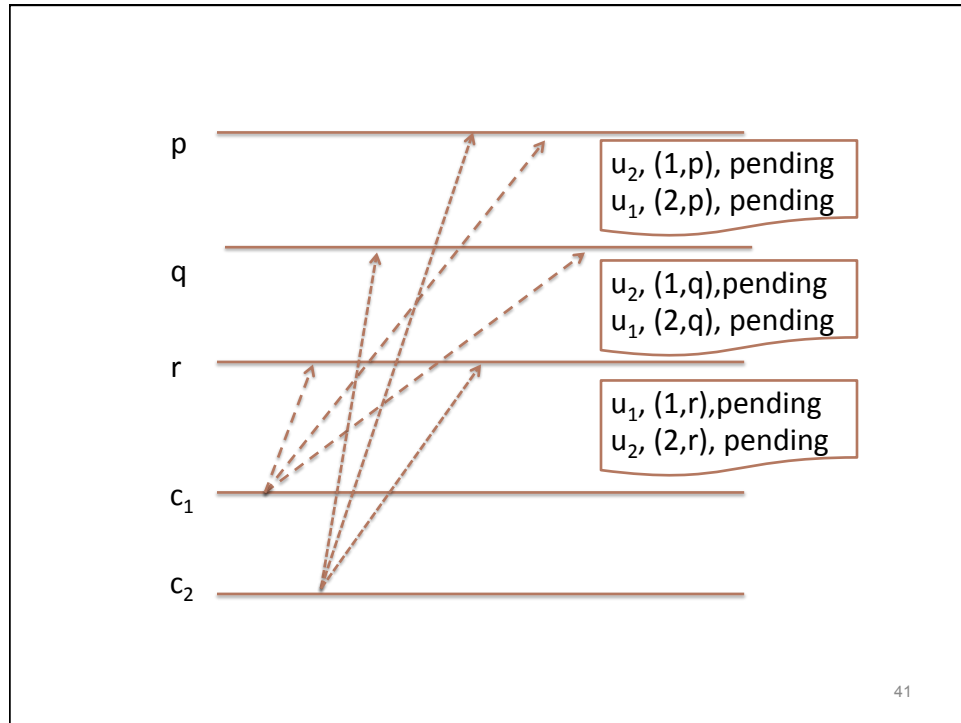
- Logical clocks
 - See mutual exclusion algorithm with logical time
- Update source takes the maximum
 - Commit message: “commit at [t,pid]”
 - Group member: if vote considered:
 - deliver committed updates in timestamp order
 - Group members: if vote not considered:
 - discard the update


39

Where are the updates?

- Each member: queue of uncommitted updates
 - Survives crash and restart
- Example: Process p
 - (u_2 : [1,p] pending), (u_1 : [2,p] pending)
 - Neither can be delivered

40





43

STATE MACHINE: PROTOCOL ANALYSIS

44

What if “my vote wasn’t used?”

- Process
 - had a pending update
 - discovers it wasn’t used
- Discard the request
 - Otherwise block forever (why?)
 - Ignoring the request won’t hurt (why?)

45

Which votes got counted?

- Need to know which votes were “counted”
 - E.g. suppose A,B,C,D,E and they vote:
 - {[17,A] [19,B] [20,C] [200,D] [21,E]}
 - Vote from D is lost
 - the maximum is picked as [21,E]
 - Remember that the votes used to make this decision were from {A,B,C,E}

46

Recovery

- First recover queue of pending updates
- Next, learn the outcome of the operation
 - Contact Q_R other replicas
- Check if own vote counted (if committed)
 - If so, apply update
 - If not, discard update

47

Read requests while updates pending...

- Suppose a read while updates pending
 - Wait until those commit, abort, or are discarded
 - Otherwise process might not see its own updated value

48

Why is this “safe”?

- Commit: only move pending update to later
 - Discard pending update if vote not counted
 - Result: inconsistent replica
 - *but we always look at Q_R replicas*
 - Why we can't support incremental operations (insert, etc)

49

Why is this “safe”?

- Commit: moves pending update to front of Q
- Once a committed update reaches front of Q:
 - ...no update can be committed at earlier time!
- Any “future” update gets later time

50

Why this works

- Everyone uses same commit time for an update
 - Can't deliver update unless [t,pid] is smallest
 - and is committed
 - Hence updates in same order at all replicas

51

Observations

- The protocol requires many messages
 - Could use IP multicast for first and last round
 - Need reliability
- Commit messages must be reliably delivered
 - Otherwise uncommitted updates on front of Q...
- 2PC and 3PC may block
 - FLP: *any* quorum write protocol can block

52

Risk of blocking

53


**PRIMARY-BACKUP:
VIEWSTAMP REPLICATION**

54

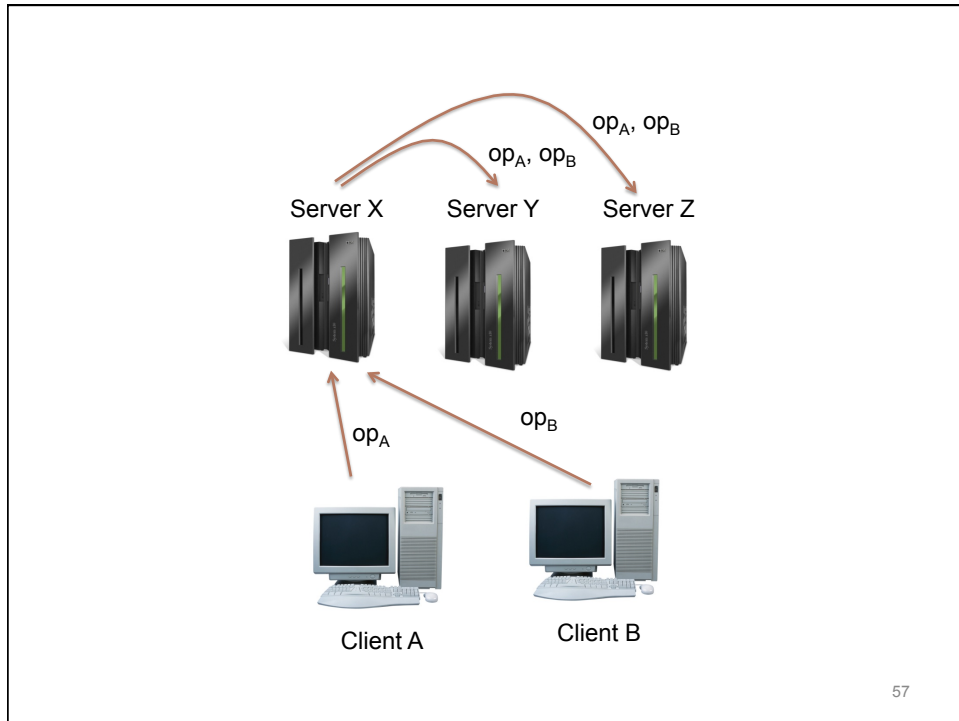
Quorum Consensus

- Crash-stop failures
- Requires $2f+1$ replicas
 - Operations must intersect for at least one replica
 - Want availability for both reads and writes
 - Read and write quorums of $f+1$ nodes

55



56



Viewstamp Replication

- Primary-backup
- System moves through a sequence of views
 - Primary runs the protocol
 - Replicas do a view change if it fails


Replica state

- A **replica id** i (between 0 and $N-1$)
 - Replica 0, replica 1, ...
- A **view number** $v\#$, initially 0
- **Primary** is the replica with id $i = v\# \bmod N$
- A **log** of $\langle op, op\#, status \rangle$ entries
 - Status = **prepared** or **committed**

59

Client knows current view #


View # = 2 \Rightarrow primary pid = 2 (i.e., r)



60


Primary chooses logical timestamp for update, e.g. $LT(u_1)=9$

Op # = (View #, time)



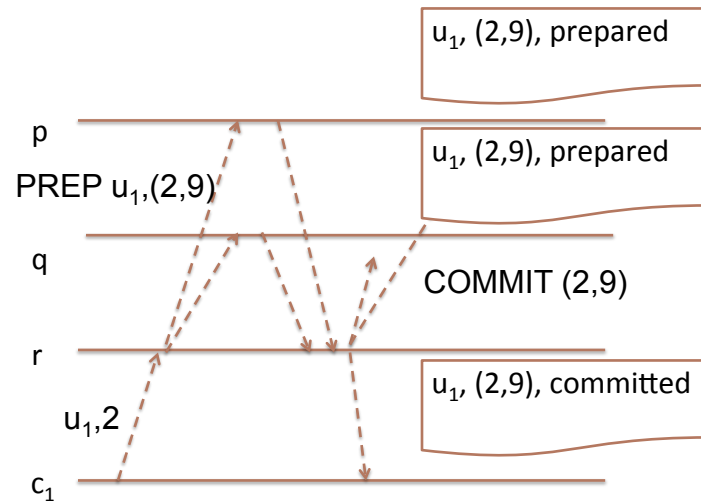
61

Primary waits for $\geq f$ replicas to respond ($\geq f+1$ total)



62

Client notified immediately after acks



63

**PRIMARY-BACKUP:
VIEW CHANGE**

64

View Changes

- Used to mask primary failures
- Replicas monitor the primary
- Replica requests next primary to do a view change


65

Correctness Requirement

- Operation order must be preserved by a view change
- For operations that are **visible**
 - executed by server
 - client received result
- An operation can be visible if it prepared at $f+1$ replicas
 - this is the **commit point**


66

Primary r has crashed after assigning $LT(u_1)=9$



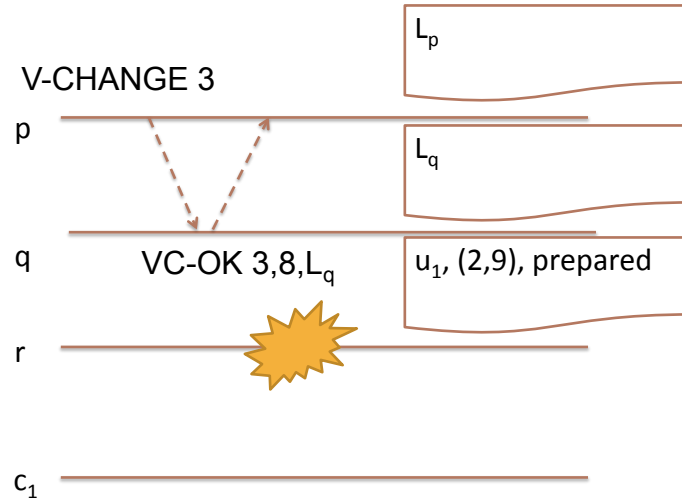
67

View # increments (to 3), new primary is p, start view change




68

Replica q acks view change, returns last op# it knew of and its log (L_q contains any operations that p never heard of)



69

Primary p lets replicas know (via its log L_p') of ops that they were never informed of before old primary crashed




70

View Change


- New primary may not know of all updates from old primary
- New primary asks replicas for their logs
- Any committed operation was acked by a quorum, so must be in log of a surviving replica
 - Primary takes the max of the logs returned
 - That log has most recent updates

Second client c_2 knows current view #

View# = 3 \Rightarrow primary pid = 0 (i.e., p)




Primary p gets the other replicas to prepare to commit u_2



73

Viewstamp (3,9) avoids confusion between u_1 and u_2 at r
(Replica r will not think p is committing u_1)



74

Persistent State

- Voting protocol: votes must survive failure
 - Save queue of pending updates on disk
- Viewstamp: primary can respond to client without recording commit on disk
 - View change: recover commitment from logs of surviving replicas
- Only need to persist state after view change
 - So if we crash and recover, we know view# when we crashed
 - Even that unnecessary with more expensive recovery protocol